

CS-C3240 – Machine Learning D

Anomaly detection, Recommender Systems, Online learning

Stephan Sigg

Department of Communications and Networking
Aalto University, School of Electrical Engineering
stephan.sigg@aalto.fi

Version 1.0, January 19, 2022

Learning goals

tbd

Outline

Recommender systems

Stochastic Classification

Online learning

Elektronik



Alle Empfehlungen in Elektronik anzeigen

Computer & Zubehör



Alle Empfehlungen in Computer & Zubehör anzeigen

Musik




Alle Empfehlungen in Musik anzeigen


Bücher




Recommender systems



Recommender systems



Recommender systems



Recommender systems












Recommender systems



Recommender systems












Recommender systems










	 ★★★★☆	 ★★★★★	 ?	 ★☆☆☆☆
	★★★★★	★★★★★	?	★☆☆☆☆
	★★★★☆	?	★☆☆☆☆	?
	?	☆☆☆☆☆	★★★★★	★★★★★
	★☆☆☆☆	★☆☆☆☆	★★★★★	?










Recommender systems

Task of Recommender systems

Given these ratings for a number of products, predict likely user-ratings for products that have not yet been rated

				
	★★★★☆	★★★★★	?	☆☆☆☆☆
	★★★★★	★★★★★	?	☆☆☆☆☆
	★★★★☆	?	☆☆☆☆☆	?
	?	☆☆☆☆☆	★★★★★	★★★★★
	☆☆☆☆☆	☆☆☆☆☆	★★★★★	?

					tools	accessories
	★★★★☆	★★★★★	?	☆☆☆☆☆	0.9	0.0
	★★★★★	★★★★★	?	★☆☆☆☆	1.0	0.3
	★★★★☆	?	★★★☆☆	?	0.4	0.8
	?	☆☆☆☆☆	★★★★★	★★★★★	0.2	0.9
	★☆☆☆☆	★☆☆☆☆	★★★★★	?	0.0	1.0

				
	★★★★☆	★★★★★	?	☆☆☆☆☆
	★★★★★	★★★★★	?	★☆☆☆☆
	★★★★☆	?	★★☆☆☆	?
	?	☆☆☆☆☆	★★★★★	★★★★★
	★☆☆☆☆	★☆☆☆☆	★★★★★	?

tools

accessories

0.9

0.0

1.0

0.3

0.4


0.8






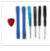



0.2










0.9

0.0

1.0












	 $W^{(1)}$	 $W^{(2)}$	 $W^{(3)}$	 $W^{(4)}$	tools	accessories
	★★★★☆	★★★★★	?	☆☆☆☆☆	0.9	$x^{(1)}$ 0.0
	★★★★★	★★★★★	?	★☆☆☆☆	1.0	$x^{(2)}$ 0.3
	★★★★☆	?	★★★☆☆	?	0.4	$x^{(3)}$ 0.8
	?	☆☆☆☆☆	★★★★★	★★★★★	0.2	$x^{(4)}$ 0.9
	★☆☆☆☆	★☆☆☆☆	★★★★★	?	0.0	$x^{(5)}$ 1.0

	 $w^{(1)}$	 $w^{(2)}$	 $w^{(3)}$	 $w^{(4)}$	tools	accessories
	★★★★☆	★★★★★	?	☆☆☆☆☆	0.9	$x^{(1)}$ 0.0
	★★★★★	★★★★★	?	★☆☆☆☆	1.0	$x^{(2)}$ 0.3
	★★★★☆	?	★★☆☆☆	?	0.4	$x^{(3)}$ 0.8
	?	☆☆☆☆☆	★★★★★	★★★★★	0.2	$x^{(4)}$ 0.9
	★☆☆☆☆	★☆☆☆☆	★★★★★	?	0.0	$x^{(5)}$ 1.0

Represent items in terms of weighted feature vectors










$$x^{(4)} = \begin{bmatrix} 0.2 \\ 0.9 \end{bmatrix} \quad w^{(1)} = \begin{bmatrix} 5 \\ 1 \end{bmatrix}$$

$$h(x) \Rightarrow \left(w^{(1)}\right)^T x^{(4)} = 0.2 \cdot 5 + 0.9 \cdot 1 = 1.9$$

	 $w^{(1)}$	 $w^{(2)}$	 $w^{(3)}$	 $w^{(4)}$	tools	accessories
	★★★★☆	★★★★★	?	☆☆☆☆☆	0.9	$x^{(1)}$ 0.0
	★★★★★	★★★★★	?	★☆☆☆☆	1.0	$x^{(2)}$ 0.3
	★★★★☆	?	★★★☆☆	?	0.4	$x^{(3)}$ 0.8
	?	☆☆☆☆☆	★★★★★	★★★★★	0.2	$x^{(4)}$ 0.9
	★☆☆☆☆	★☆☆☆☆	★★★★★	?	0.0	$x^{(5)}$ 1.0

Learn weights from provided ratings for single user j (Logistic regression):

$$\min_{w^{(j)}} \frac{1}{2} \sum_{i: y^{(i,j)} \neq ?} \left(\left(w^{(j)} \right)^T x^{(i)} - y^{(i,j)} \right)^2 + \frac{\lambda}{2} \sum_{k=1}^m \left(w_k^{(j)} \right)^2$$

	 $w^{(1)}$	 $w^{(2)}$	 $w^{(3)}$	 $w^{(4)}$	tools	accessories
	★★★★☆	★★★★★	?	☆☆☆☆☆	0.9	$x^{(1)}$ 0.0
	★★★★★	★★★★★	?	★☆☆☆☆	1.0	$x^{(2)}$ 0.3
	★★★★☆	?	★★★☆☆	?	0.4	$x^{(3)}$ 0.8
	?	☆☆☆☆☆	★★★★★	★★★★★	0.2	$x^{(4)}$ 0.9
	★☆☆☆☆	★☆☆☆☆	★★★★★	?	0.0	$x^{(5)}$ 1.0

Learn weights from provided ratings for all users $1, \dots, N$:

$$\min_{w^{(1)}, \dots, w^{(N)}} \frac{1}{2} \sum_{j=1}^N \sum_{i: y^{(i,j)} \neq ?} \left((w^{(j)})^T x^{(i)} - y^{(i,j)} \right)^2 + \frac{\lambda}{2} \sum_{j=1}^N \sum_{k=1}^m (w_k^{(j)})^2$$

Recommender systems

Optimisation algorithm

$$\min_{w^{(1)}, \dots, w^{(N)}} \frac{1}{2} \sum_{j=1}^N \sum_{i: y^{(i,j)} \neq ?} \left(\left(w^{(j)} \right)^T x^{(i)} - y^{(i,j)} \right)^2 + \frac{\lambda}{2} \sum_{j=1}^N \sum_{k=1}^m \left(w_k^{(j)} \right)^2$$

Recommender systems

Optimisation algorithm

$$\min_{w^{(1)}, \dots, w^{(N)}} \frac{1}{2} \sum_{j=1}^N \sum_{i: y^{(i,j)} \neq ?} \left(\left(w^{(j)} \right)^T x^{(i)} - y^{(i,j)} \right)^2 + \frac{\lambda}{2} \sum_{j=1}^N \sum_{k=1}^m \left(w_k^{(j)} \right)^2$$

Gradient descent update:

$$w_k^{(j)} = w_k^{(j)} - \delta \left(\sum_{i: y^{(i,j)} \neq ?} \left(\left(w^{(j)} \right)^T x^{(i)} - y^{(i,j)} \right) x_k^{(i)} + \lambda w_k^{(j)} \right)$$

Recommender systems

Optimisation algorithm










$$\min_{w^{(1)}, \dots, w^{(N)}} \frac{1}{2} \sum_{j=1}^N \sum_{i: y^{(i,j)} \neq ?} \left(\left(w^{(j)} \right)^T x^{(i)} - y^{(i,j)} \right)^2 + \frac{\lambda}{2} \sum_{j=1}^N \sum_{k=1}^m \left(w_k^{(j)} \right)^2$$

Gradient descent update:

$$w_k^{(j)} = w_k^{(j)} - \delta \left(\underbrace{\sum_{i: y^{(i,j)} \neq ?} \left(\left(w^{(j)} \right)^T x^{(i)} - y^{(i,j)} \right) x_k^{(i)}}_{\text{partial derivative}} + \lambda w_k^{(j)} \right)$$

Recommender systems

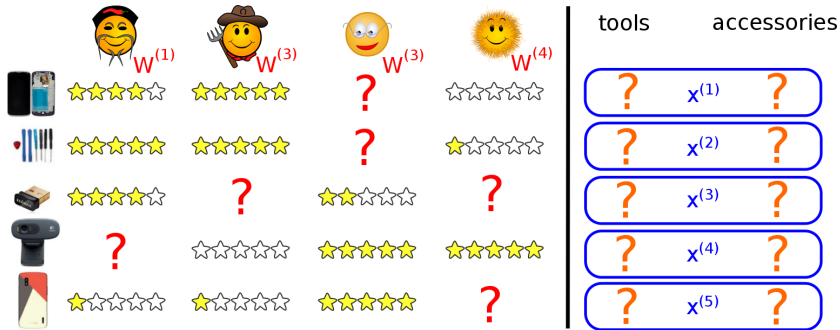
Collaborative filtering

	 $w^{(1)}$	 $w^{(2)}$	 $w^{(3)}$	 $w^{(4)}$	tools	accessories
	★★★★☆	★★★★★	?	☆☆☆☆☆	0.9	$x^{(1)}$ 0.0
	★★★★★	★★★★★	?	★☆☆☆☆	1.0	$x^{(2)}$ 0.3
	★★★★☆	?	★★☆☆☆	?	0.4	$x^{(3)}$ 0.8
	?	☆☆☆☆☆	★★★★★	★★★★★	0.2	$x^{(4)}$ 0.9
	★☆☆☆☆	★☆☆☆☆	★★★★★	?	0.0	$x^{(5)}$ 1.0

We are able to calculate the weights given the feature vectors

Recommender systems

Collaborative filtering



We are able to calculate the weights given the feature vectors








→ But how do we obtain these feature vectors?

Recommender systems

Collaborative filtering

Users might tell their preferences

e.g. more interested in tools or accessories

	tools	accessories
	 $w^{(1)}$ ★★★★★	 $w^{(3)}$ ★★★★★
	★★★★★	★★★★★
	★★★★★	?
	?	★★★★★
	★★★★★	?








$$w^{(1)} = \begin{bmatrix} 5 \\ 0 \end{bmatrix} \quad w^{(2)} = \begin{bmatrix} 5 \\ 0 \end{bmatrix} \quad w^{(3)} = \begin{bmatrix} 0 \\ 5 \end{bmatrix} \quad w^{(4)} = \begin{bmatrix} 0 \\ 5 \end{bmatrix}$$

Recommender systems

Collaborative filtering

Users might tell their preferences

e.g. more interested in tools or accessories

	tools	accessories
	 $w^{(1)}$ ★★★★☆	 $w^{(2)}$ ★★★★★
	★★★★★	★★★★★
	★★★★★	?
	?	★★★★★
	★★★★★	★★★★★

$$w^{(1)} = \begin{bmatrix} 5 \\ 0 \end{bmatrix} \quad w^{(2)} = \begin{bmatrix} 5 \\ 0 \end{bmatrix} \quad w^{(3)} = \begin{bmatrix} 0 \\ 5 \end{bmatrix} \quad w^{(4)} = \begin{bmatrix} 0 \\ 5 \end{bmatrix}$$

$$\left(w^{(1)}\right)^T x^{(1)} \approx 4; \quad \left(w^{(2)}\right)^T x^{(1)} \approx 5;$$








$$\left(w^{(3)}\right)^T x^{(1)} \approx ?; \quad \left(w^{(4)}\right)^T x^{(1)} \approx 0$$

Recommender systems

Collaborative filtering

Users might tell their preferences

e.g. more interested in tools or accessories

	tools	accessories
	 $w^{(1)}$ ★★★★☆	 $w^{(2)}$ ★★★★★
	★★★★★	★★★★★
	★★★★★	?
	?	★★★★★
	★★★★★	★★★★★

$$w^{(1)} = \begin{bmatrix} 5 \\ 0 \end{bmatrix} \quad w^{(2)} = \begin{bmatrix} 5 \\ 0 \end{bmatrix} \quad w^{(3)} = \begin{bmatrix} 0 \\ 5 \end{bmatrix} \quad w^{(4)} = \begin{bmatrix} 0 \\ 5 \end{bmatrix}$$

$$\left(w^{(1)}\right)^T x^{(1)} \approx 4; \quad \left(w^{(2)}\right)^T x^{(1)} \approx 5;$$

$$\left(w^{(3)}\right)^T x^{(1)} \approx ?; \quad \left(w^{(4)}\right)^T x^{(1)} \approx 0$$

From these weights we can estimate the feature values

Recommender systems

Collaborative filtering

Optimisation algorithm

Given the weights/preferences $w^{(1)}, \dots, w^{(N)}$, we are able to infer a feature $x^{(i)}$

$$\min_{x^{(i)}} \frac{1}{2} \sum_{j: y^{(i,j)} \neq ?} \left(\left(w^{(j)} \right)^T x^{(i)} - y^{(i,j)} \right)^2 + \frac{\lambda}{2} \sum_{k=1}^m \left(x_k^{(i)} \right)^2$$

Recommender systems

Collaborative filtering

Optimisation algorithm

Given the weights/preferences $w^{(1)}, \dots, w^{(N)}$, we are able to infer a feature $x^{(i)}$

$$\min_{x^{(i)}} \frac{1}{2} \sum_{j: y^{(i,j)} \neq ?} \left((w^{(j)})^T x^{(i)} - y^{(i,j)} \right)^2 + \frac{\lambda}{2} \sum_{k=1}^m (x_k^{(i)})^2$$

Given the weights/preferences $w^{(1)}, \dots, w^{(N)}$, we are able to infer $x^{(1)}, \dots, x^{(n)}$

$$\min_{x^{(1)}, \dots, x^{(n)}} \frac{1}{2} \sum_{i=1}^n \sum_{j: y^{(i,j)} \neq ?} \left((w^{(j)})^T x^{(i)} - y^{(i,j)} \right)^2 + \frac{\lambda}{2} \sum_{i=1}^n \sum_{k=1}^m (x_k^{(i)})^2$$

Recommender systems

Collaborative filtering

Naive (iterative) Collaborative filtering algorithm

Given $x^{(1)}, \dots, x^{(n)}$, we are able to estimate $w^{(1)}, \dots, w^{(N)}$

Given $w^{(1)}, \dots, w^{(N)}$, we are able to estimate $x^{(1)}, \dots, x^{(n)}$

Recommender systems

Collaborative filtering

Naive (iterative) Collaborative filtering algorithm

Given $x^{(1)}, \dots, x^{(n)}$, we are able to estimate $w^{(1)}, \dots, w^{(N)}$

Given $w^{(1)}, \dots, w^{(N)}$, we are able to estimate $x^{(1)}, \dots, x^{(n)}$

Collaborative filtering (naive)

Init: Randomly initialise the $w^{(i)}$

Repeat:

- Estimate the $x^{(i)}$ from the $w^{(i)}$
- Estimate the $w^{(i)}$ from the $x^{(i)}$

Recommender systems

Collaborative filtering

Naive (iterative) Collaborative filtering algorithm

Given $x^{(1)}, \dots, x^{(n)}$, we are able to estimate $w^{(1)}, \dots, w^{(N)}$

Given $w^{(1)}, \dots, w^{(N)}$, we are able to estimate $x^{(1)}, \dots, x^{(n)}$

Collaborative filtering (naive)

Init: Randomly initialise the $w^{(i)}$

Repeat:

- Estimate the $x^{(i)}$ from the $w^{(i)}$
- Estimate the $w^{(i)}$ from the $x^{(i)}$

- CF iteratively improves the estimates for $x^{(i)}$ and $w^{(i)}$
- Algorithm collaborates with users: by providing some information about their preferences, it computes and improves the features

Recommender systems

Collaborative filtering

$$\min_{w^{(1)}, \dots, w^{(N)}} \quad \frac{1}{2} \sum_{j=1}^N \sum_{i: y^{(i,j)} \neq ?} \left((w^{(j)})^T x^{(i)} - y^{(i,j)} \right)^2 + \frac{\lambda}{2} \sum_{j=1}^N \sum_{k=1}^m \left(w_k^{(j)} \right)^2$$
$$\min_{x^{(1)}, \dots, x^{(n)}} \quad \frac{1}{2} \sum_{i=1}^n \sum_{j: y^{(i,j)} \neq ?} \left((w^{(j)})^T x^{(i)} - y^{(i,j)} \right)^2 + \frac{\lambda}{2} \sum_{i=1}^n \sum_{k=1}^m \left(x_k^{(i)} \right)^2$$

Recommender systems

Collaborative filtering

$$\min_{w^{(1)}, \dots, w^{(N)}} \quad \frac{1}{2} \sum_{j=1}^N \sum_{i: y^{(i,j)} \neq ?} \left((w^{(j)})^T x^{(i)} - y^{(i,j)} \right)^2 + \frac{\lambda}{2} \sum_{j=1}^N \sum_{k=1}^m \left(w_k^{(j)} \right)^2$$
$$\min_{x^{(1)}, \dots, x^{(n)}} \quad \frac{1}{2} \sum_{i=1}^n \sum_{j: y^{(i,j)} \neq ?} \left((w^{(j)})^T x^{(i)} - y^{(i,j)} \right)^2 + \frac{\lambda}{2} \sum_{i=1}^n \sum_{k=1}^m \left(x_k^{(i)} \right)^2$$

Recommender systems

Collaborative filtering

$$\min_{w^{(1)}, \dots, w^{(N)}} \quad \frac{1}{2} \sum_{j=1}^N \sum_{i: y^{(i,j)} \neq ?} \left((w^{(j)})^T x^{(i)} - y^{(i,j)} \right)^2 + \frac{\lambda}{2} \sum_{j=1}^N \sum_{k=1}^m \left(w_k^{(j)} \right)^2$$

$$\min_{x^{(1)}, \dots, x^{(n)}} \quad \frac{1}{2} \sum_{i=1}^n \sum_{j: y^{(i,j)} \neq ?} \left((w^{(j)})^T x^{(i)} - y^{(i,j)} \right)^2 + \frac{\lambda}{2} \sum_{i=1}^n \sum_{k=1}^m \left(x_k^{(i)} \right)^2$$

Minimize $w^{(i)}$ and $x^{(i)}$ simultaneously:

$$\begin{aligned} \min_{x^{(1)}, \dots, x^{(n)}, w^{(1)}, \dots, w^{(N)}} \quad & \frac{1}{2} \sum_{i,j: y^{(i,j)} \neq ?} \left((w^{(j)})^T x^{(i)} - y^{(i,j)} \right)^2 \\ & + \frac{\lambda}{2} \sum_{i=1}^n \sum_{k=1}^m \left(x_k^{(i)} \right)^2 + \frac{\lambda}{2} \sum_{j=1}^N \sum_{k=1}^m \left(w_k^{(j)} \right)^2 \end{aligned}$$

Recommender systems

Init: Randomly initialise the $w^{(j)}$ and $x^{(i)}$

Optimisation: Simultaneously minimise the above function for $w^{(j)}$ and $x^{(i)}$

Gradient descent:

$$\begin{aligned}x_k^{(i)} &= x_k^{(i)} - \delta \left(\sum_{j=y^{(i,j)} \neq ?} \left((w^{(j)})^T x^{(i)} - y^{(i,j)} \right) w_k^{(j)} + \lambda x_k^{(i)} \right) \\w_k^{(j)} &= w_k^{(j)} - \delta \left(\sum_{i=y^{(i,j)} \neq ?} \left((w^{(j)})^T x^{(i)} - y^{(i,j)} \right) x_k^{(i)} + \lambda w_k^{(j)} \right)\end{aligned}$$

Prediction: For a user i with parameters $w^{(j)}$ and an item with learned features x , estimate a rating of $(w^{(j)})^T x$

Outline

Recommender systems

Stochastic Classification

Online learning

Stochastic Classification

Model training becomes slow with increasing data size

Stochastic Classification

Model training becomes slow with increasing data size

- Because of repeated looping over the complete data set until convergence

Stochastic Classification

Model training becomes slow with increasing data size


→ Because of repeated looping over the complete data set until convergence

Solution

- Randomly iterate the update only over a subset of items instead of repeatedly considering the whole data set.

Stochastic Classification

Example: Gradient descent




$$\text{minimize } L[(\mathcal{X}, \mathcal{Y}), h(\cdot)] = \frac{1}{2n} \sum_{i=1}^n \left(w^T x^{(i)} - y^{(i)} \right)^2$$

$$\text{Repeat } \forall j : w_j = w_j - \delta \cdot \frac{\partial L[(\mathcal{X}, \mathcal{Y}), h(\cdot)]}{\partial w_j}$$

$$\rightarrow \forall j : w_j = w_j - \delta \cdot \frac{1}{n} \sum_{i=1}^n \left(w_j x_j^{(i)} - y^{(i)} \right) \cdot x_j^{(i)}$$

Stochastic Classification

Example: Gradient descent



→ For single gradient descent-step, algorithms loops over all samples!

$$\text{minimize } L[(\mathcal{X}, \mathcal{Y}), h(\cdot)] = \frac{1}{2n} \sum_{i=1}^n \left(w^T x^{(i)} - y^{(i)} \right)^2$$

$$\text{Repeat } \forall j : w_j = w_j - \delta \cdot \frac{\partial L[(\mathcal{X}, \mathcal{Y}), h(\cdot)]}{\partial w_j}$$

$$\rightarrow \forall j : w_j = w_j - \delta \cdot \frac{1}{n} \sum_{i=1}^n \left(w_j x_j^{(i)} - y^{(i)} \right) \cdot x_j^{(i)}$$

Stochastic Classification

Example: Gradient descent

→ For a single gradient descent-step, the algorithms loops over all samples!

$$\text{minimize } L[(\mathcal{X}, \mathcal{Y}), h(\cdot)] = \frac{1}{2n} \sum_{i=1}^n \left(w^T x^{(i)} - y^{(i)} \right)^2$$

$$\text{Repeat } \forall j : w_j = w_j - \delta \cdot \frac{\partial L[(\mathcal{X}, \mathcal{Y}), h(\cdot)]}{\partial w_j}$$

$$\rightarrow \forall j : w_j = w_j - \delta \cdot \frac{1}{n} \sum_{i=1}^n \left(w_j x_j^{(i)} - y^{(i)} \right) \cdot x_j^{(i)}$$

Stochastic Classification

Example: Gradient descent

→ For a single gradient descent-step, the algorithm loops over all samples!

To speed up the algorithm, compute gradient descent updates from individual training samples (randomly ordered)

$$\text{minimize } L[(\mathcal{X}, \mathcal{Y}), h(\cdot)] = \frac{1}{2n} \sum_{i=1}^n \left(w^T x^{(i)} - y^{(i)} \right)^2$$

$$\text{Repeat } \forall j : w_j = w_j - \delta \cdot \frac{\partial L[(\mathcal{X}, \mathcal{Y}), h(\cdot)]}{\partial w_j}$$

$$\rightarrow \forall j : w_j = w_j - \delta \cdot \frac{1}{n} \sum_{i=1}^n \left(w_j x_j^{(i)} - y^{(i)} \right) \cdot x_j^{(i)}$$

Stochastic Classification

Example: Gradient descent

Standard:

$$\text{minimize } L[(\mathcal{X}, \mathcal{Y}), h(\cdot)] = \frac{1}{2n} \sum_{i=1}^n \left(w^T x^{(i)} - y^{(i)} \right)^2$$

$$\text{Repeat } \forall j : w_j = w_j - \delta \cdot \frac{\partial}{\partial w_j} E[w_j]$$

$$\rightarrow \forall j : w_j = w_j - \delta \cdot \frac{1}{n} \sum_{i=1}^n \left(w_j x_j^{(i)} - y^{(i)} \right) \cdot x_j^{(i)}$$

Stochastic Classification

Example: Gradient descent

Standard:

$$\text{minimize } E[W] = \frac{1}{2n} \sum_{i=1}^n \left(w^T x^{(i)} - y^{(i)} \right)^2$$

$$\text{Repeat } \forall j : w_j = w_j - \delta \cdot \frac{\partial}{\partial w_j} E[w_j]$$

$$\rightarrow \forall j : w_j = w_j - \delta \cdot \frac{1}{n} \sum_{i=1}^n \left(w_j x_j^{(i)} - y^{(i)} \right) \cdot x_j^{(i)}$$

Stochastic:

Repeat over all training examples i (random order):

$$\Rightarrow \forall j : w_j = w_j - \delta \cdot \left(w_j x_j^{(i)} - y^{(i)} \right) \cdot x_j^{(i)}$$

Stochastic Classification

The stochastic implementation will generally move the parameters towards the global minimum (...but not always!)

Stochastic Classification

The stochastic implementation will generally move the parameters towards the global minimum (...but not always!)

Greatly speeds up the gradient descent steps as it does not loop over all samples in each single iteration

Stochastic Classification

The stochastic implementation will generally move the parameters towards the global minimum (...but not always!)

Greatly speeds up the gradient descent steps as it does not loop over all samples in each single iteration

Tradeoff use $1 \leq k \leq n$ random examples for each gradient descent update

$l = 1, 1 + k, 1 + 2k, \dots$

$$\Rightarrow \forall j : w_j = w_j - \delta \cdot \frac{1}{l} \sum_{l=i}^{i+k} \left(w_j x_j^{(i)} - y^{(i)} \right) \cdot x_j^{(i)}$$

Stochastic Classification

The stochastic implementation will generally move the parameters towards the global minimum (...but not always!)

Greatly speeds up the gradient descent steps as it does not loop over all samples in each single iteration

Tradeoff use $1 \leq k \leq n$ random examples for each gradient descent update

$l = 1, 1 + k, 1 + 2k, \dots$

$$\Rightarrow \forall j : w_j = w_j - \delta \cdot \frac{1}{l} \sum_{l=i}^{i+k} \left(w_j x_j^{(i)} - y^{(i)} \right) \cdot x_j^{(i)}$$

$k > 1$ might be faster than $k = 1$ for parallelized code

Outline

Recommender systems

Stochastic Classification

Online learning

Online learning

Problem formulation

Online learning

Given a continuous stream of input data, update the parameters of your algorithm on-the-fly

Example: learn user behaviour from website users

Online learning

Problem formulation

Online learning

Given a continuous stream of input data, update the parameters of your algorithm on-the-fly

Example: learn user behaviour from website users

Similar to stochastic classification: Update the parameters based on individual training examples

$$\Rightarrow \forall j : w_j = w_j - \delta \cdot (h(x^{(i)}) - y^{(i)}) \cdot x_j^{(i)}$$

Online learning

Problem formulation

Online learning

Given a continuous stream of input data, update the parameters of your algorithm on-the-fly

Example: learn user behaviour from website users

Similar to stochastic classification: Update the parameters based on individual training examples

$$\Rightarrow \forall j : w_j = w_j - \delta \cdot (h(x^{(i)}) - y^{(i)}) \cdot x_j^{(i)}$$

\Rightarrow Able to adapt to changing user behaviour over time

Questions?

Stephan Sigg

`stephan.sigg@aalto.fi`

Si Zuo

`si.zuo@aalto.fi`

Literature

- C.M. Bishop: Pattern recognition and machine learning, Springer, 2007.
- R.O. Duda, P.E. Hart, D.G. Stork: Pattern Classification, Wiley, 2001.

